

# Edge detection optimization

December 17, 2023

Adrien Rey

## 1 Introduction

For the Embedded Hardware project, we have to optimize the edge detection on a Nios II. A working project was given but the performance are not good. The given project has a performance of 8450 cycles per pixels. By following the theory seen during the course, the goal is to improve the performance of the system.

## 2 Software optimization

### 2.1 Compiler optimization

By modifying the compiler optimization, the performance is improved without changing the code :

🌀 With -O0 : 8450 cycles/pixel

🌀 With -O1 : 1570 cycles/pixel

🌀 With -O2 : 1337 cycles/pixel

🌀 With -O3 : 545 cycles/pixel

### 2.2 Sobel

#### 2.2.1 Loop unrolling

By unrolling the inner and outer loop of the *sobel\_mac* function the performance is improve. The loop was unroll 9 times so it no need loop anymore. The function has the following form.

```
short sobel_mac( unsigned char *pixels,
                 int y,
                 const char *filter,
                 unsigned int width ) {
    short dy,dx;
    short result = 0;

    //dy = -1
    result += filter[0]*
        pixels[(y-1)*width+(x-1)];
    result += filter[1]*
        pixels[(y-1)*width+x];
    result += filter[2]*
        pixels[(y-1)*width+(x+1)];
    //dy = 0
    result += filter[3]*
        pixels[y*width+(x-1)];
    result += filter[4]*
        pixels[y*width+x];
    result += filter[5]*
        pixels[y*width+(x+1)];
    //dy = 1
    result += filter[6]*
        pixels[(y+1)*width+(x-1)];
    result += filter[7]*
        pixels[(y+1)*width+x];
```

```

    result += filter[8]*
        pixels[(y+1)*width+(x+1)];

    return result;
}

```

🌀 With -O0 : 5606 cycles/pixel

### 2.2.2 In-lining

The function *sobel\_mac* is integrated directly in the *sobel\_x* and the *sobel\_y* function as shown in the code snippet below. With this optimization, the performance is much improved.

```

void sobel_x(unsigned char *source) {
    int x, y;

    for (y = 1; y < (sobel_height - 1); y++) {
        for (x = 1; x < (sobel_width - 1); x++) {
            sobel_x_result[y * sobel_width + x] =
                currentRow[x - 1 - sobel_width] * gx_array[0][0] +
                currentRow[x - sobel_width] * gx_array[0][1] +
                currentRow[x + 1 - sobel_width] * gx_array[0][2] +
                currentRow[x - 1] * gx_array[1][0] +
                currentRow[x] * gx_array[1][1] +
                currentRow[x + 1] * gx_array[1][2] +
                currentRow[x - 1 + sobel_width] * gx_array[2][0] +
                currentRow[x + sobel_width] * gx_array[2][1] +
                currentRow[x + 1 + sobel_width] * gx_array[2][2];
        }
    }
}

```

🌀 With -O0 : 3855 cycles/pixel

🌀 With -O3 : 630 cycles/pixel

### 2.2.3 Pointer

To further improve the performance, pointers were used to access datas and write new values as in the following code.

```

void sobel_x(unsigned char *source) {
    int x, y;

    for (y = 1; y < (sobel_height - 1); y++) {
        for (x = 1; x < (sobel_width - 1); x++) {
            unsigned char *currentRow = source + y * sobel_width;

            currentRow[y * sobel_width + x] =
                currentRow[x - 1 - sobel_width] * gx_array[0][0] +
                currentRow[x - sobel_width] * gx_array[0][1] +
                currentRow[x + 1 - sobel_width] * gx_array[0][2] +
                currentRow[x - 1] * gx_array[1][0] +
                currentRow[x] * gx_array[1][1] +
                currentRow[x + 1] * gx_array[1][2] +
                currentRow[x - 1 + sobel_width] * gx_array[2][0] +
                currentRow[x + sobel_width] * gx_array[2][1] +
                currentRow[x + 1 + sobel_width] * gx_array[2][2];
        }
    }
}

```

```

    }
}

```

🔗 With -O3 : 545 cycles/pixel

### 2.2.4 Complete method

All the logic for the sobel was reunited a method called *sobel*. The contain of the *sobel\_x* and *sobel\_y* are embedded in this function. The code from the threshold is also included. With this, the double for loops from the *sobel\_threshold*, *sobel\_x* and *sobel\_y* are reunited in one loop. The performance are better as the number of loops decrease. In this part, local variable are use to store the x and y value. So the array is acces only one time at the end for both x and y. The threshold is simplified by the use of abs. It do not change the perfomance but it improved the code lisibility

```

void sobel(unsigned char *source, short threshold) {
    int x, y;

    for (y = 1; y < (sobel_height - 1); y++) {
        for (x = 1; x < (sobel_width - 1); x++) {
            int arrayIndex = y * sobel_width + x;
            unsigned char *currentRow = source + y * sobel_width;

            short gx = currentRow[x - 1 - sobel_width] * gx_array[0][0] +
                currentRow[x - sobel_width] * gx_array[0][1] +
                currentRow[x + 1 - sobel_width] * gx_array[0][2] +
                currentRow[x - 1] * gx_array[1][0] +
                currentRow[x] * gx_array[1][1] +
                currentRow[x + 1] * gx_array[1][2] +
                currentRow[x - 1 + sobel_width] * gx_array[2][0] +
                currentRow[x + sobel_width] * gx_array[2][1] +
                currentRow[x + 1 + sobel_width] * gx_array[2][2];

            short gy = currentRow[x - 1 - sobel_width] * gy_array[0][0] +
                currentRow[x - sobel_width] * gy_array[0][1] +
                currentRow[x + 1 - sobel_width] * gy_array[0][2] +
                currentRow[x - 1] * gy_array[1][0] +
                currentRow[x] * gy_array[1][1] +
                currentRow[x + 1] * gy_array[1][2] +
                currentRow[x - 1 + sobel_width] * gy_array[2][0] +
                currentRow[x + sobel_width] * gy_array[2][1] +
                currentRow[x + 1 + sobel_width] * gy_array[2][2];

            short sum = abs(gx) + abs(gy);

            sobel_result[arrayIndex] = (sum > threshold) ? 0xFF : 0;
        }
    }
}

```

🔗 With -O3 : 375 cycles/pixel

## 2.3 Grayscale

The grayscale function has also be improved by simplifying the equation. It has be reunited in a single line. The division by 100 has been replaced by a right shift of 128. To keep the same value at the end, the color weights have been updated following this rule  $new\_value = \frac{128}{100} * value$ .


```

void conv_grayscale(void *picture,
                    int width,
                    int height) {
    int x,y,gray;
    unsigned short *pixels = (unsigned short *)picture , rgb;
    grayscale_width = width;
    grayscale_height = height;
    if (grayscale_array != NULL)
        free(grayscale_array);
    grayscale_array = (unsigned char *) malloc(width*height);
    for (y = 0 ; y < height ; y++) {
        for (x = 0 ; x < width ; x++) {
            rgb = pixels[y*width+x];

            gray = (((rgb>>11)&0x1F)<<3)*27 +
                    (((rgb>>5)&0x3F)<<2)*92 +
                    (((rgb>>0)&0x1F)<<3)*9)
                    >> 7;


            IOWR_8DIRECT(grayscale_array,y*width+x,gray);
        }
    }
}

```

 With -O3 : 334 cycles/pixel

### 3 Cache

To optimize the memory use, the data and instruction caches have been activated with a size of 16k-bytes each. It made a huge difference to activate the cache.

 With -O3 : 140 cycles/pixel

### 4 Custom instructions

A custom instruction was written and used for the threshold. It allow to make the threshold calcul in a only cycle. The vhdl code is the following.

```

library ieee;
use ieee.std_logic_1164.all;
use ieee.std_logic_unsigned.all;
use ieee.std_logic_arith.all;

entity sobel_threshold_ci is
    port (
        signal sum          : in std_logic_vector(31 downto 0);
        signal threshold     : in std_logic_vector(31 downto 0);
        signal thresholdResult : out std_logic_vector(31 downto 0)
    );
end sobel_threshold_ci;

architecture behavioral of sobel_threshold_ci is
begin
    process(sum)
    begin
        if sum > threshold then
            thresholdResult <= (others=>'1');
        else

```

```

        thresholdResult <= (others=>'0');
    end if;
end process;
end behavioral;

```

In the *sobel* function, this line of code


```
sobel_result[arrayIndex] = (sum > threshold) ? 0xFF : 0;
```

is replaced by this one using the macro created along with the new component based on the vhdl added in Qsys

```

sobel_result[arrayIndex] = (sum > threshold) ? 0xFF : 0;
sobel_result[arrayIndex] = ALT_CI_THRESHOLD_CI_0(sum, threshold);

```

 With -O3 : 134 cycles/pixel

## 5 Memory access

The last improvements were to optimize the use of the cache memory. To do so, in the *sobel* function, instead of using a pointer to access the data, we made local copy of each of the three lines used for the calcul. We are using a *memcpy* to copy with the cache the content of the array into the local array.

```

void sobel(unsigned char *source, short threshold) {
    int x, y;

    for (y = 1; y < (sobel_height - 1); y++) {
        unsigned char l1[sobel_width];
        unsigned char l2[sobel_width];
        unsigned char l3[sobel_width];

        memcpy(l1, source + (y - 1) * sobel_width, sobel_width);
        memcpy(l2, source + y * sobel_width, sobel_width);
        memcpy(l3, source + (y + 1) * sobel_width, sobel_width);

        for (x = 1; x < (sobel_width - 1); x++) {
            int arrayIndex = y * sobel_width + x;

            short gx = l1[x - 1] * gx_array[0][0] +
                l1[x] * gx_array[0][1] +
                l1[x + 1] * gx_array[0][2] +
                l2[x - 1] * gx_array[1][0] +
                l2[x] * gx_array[1][1] +
                l2[x + 1] * gx_array[1][2] +
                l3[x - 1] * gx_array[2][0] +
                l3[x] * gx_array[2][1] +
                l3[x + 1] * gx_array[2][2];

            short gy = l1[x - 1] * gy_array[0][0] +
                l1[x] * gy_array[0][1] +
                l1[x + 1] * gy_array[0][2] +
                l2[x - 1] * gy_array[1][0] +
                l2[x] * gy_array[1][1] +
                l2[x + 1] * gy_array[1][2] +
                l3[x - 1] * gy_array[2][0] +
                l3[x] * gy_array[2][1] +
                l3[x + 1] * gy_array[2][2];

```

```

        l3[x + 1] * gy_array[2][2];

        short sum = abs(gx) + abs(gy);

        sobel_result[arrayIndex] = ALT_CI_THRESHOLD_CI_0(sum, threshold);
    }
}
}

```

🚀 With -O3 : 103 cycles/pixel

## 6 Pipeline

We have activate the pipeline in Qsys as shown in the following picture.

System: base\_system Path: nios2\_gen2\_0

Nios II Processor  
altera\_nios2\_gen2 Details

Main Vectors Caches and Memory Interfaces Arithmetic Instructions MMU and MPU Settings JTAG Debug Advanced Features

Multiply/Shift/Rotate Hardware:

Divide Hardware:

**Arithmetic Implementation**

Multiply Implementation:

Multiply Extended Implementation:

Shift/Rotate Implementation:

**Summary**

Operation	Performance	Resources	Instructions
Multiply	1 cycle	3 16-bit multipliers	MUL, MULI
Multiply Extended	2 cycles	1 extra 16-bit multiplier	MULXSS, MULXSU, MULXUU
Shift/rotate	1 cycle	Logic elements (pipelined)	ROL, ROLI, ROR, SLL, SLLI, SRA, SRAI, SRL, SRLI
Divide	35 cycles	Logic elements	DIV, DIVU

Figure 1: Settings Pipeline

🚀 With -O3 : 92 cycles/pixel

## 7 Conclusion

The final performance of the system are the following :

🚀 With -O3 : 97 cycles/pixel (31 for grayscale and 66 for the sobel)

🚀 With -O3 : 2.59 FPS

This performance are a mean of 100 measures.

Other improvements could have been done as custom instructions for the grayscale or sobel.